

Input

prediction
ground_truth
answer_format

*one row from
the eval trace*

Parse

per-format deterministic rule

canonical forms:

A · TFTF · BDCA
0.42 · 0.5_-1.2_0.0

emits (score, provenance)

Judge Extract

*LLM transcribes,
never grades*

“the final answer is
B.”

→ B

*returns canonical str. or
NO_ANSWER*

Re-parse

same Step 1 parser

B vs. ground_truth
→ score ✓

*provenance=judge
(logged for audit)*

Output

score ∈
[0, 1]
+ provenance

*parser
judge
unparseable*

parsed

NO_ANSWER ⇒ score = 0